

Coursework One - Q3 - Wisconsin Cancer - Student

Please List Members of Team and Student IDs below:

Member 1: Marcel Santos de Carvalho 79083

Member 2: Loris Baudry 79794

Member 3: Alejandro Palacios García 73713

Which member of the team is responsible for this notebook: Alejandro Palacios García

Instructions

This is a well-known dataset on predicting cancer. It involves the analysis of tumours. A tumour can be benign (not cancerous) or malignant (cancerous). Using various geometrical measurements of actual tumours, the objective is to determine whether or not a tumour is benign or malignant.

The purpose of this question is to find the best model for predicting malignancy correctly. The data is provided by Sklearn using

```
from sklearn.datasets import load_breast_cancer
```

Some of the data fields are differently scaled values. In the following you are to

1. Load the data from Scikit learn using

```
from sklearn.datasets import load_breast_cancer
```

```
data = load_breast_cancer()
```

- 1. Examine the dataframe for missing data and decide how to treat these**
- 2. Try to see which features have the best explanatory power.**
- 3. Also, deal with categorical data and use feature standardization**
- 4. Choose metrics for model evaluation**
- 5. Use the following models - Logistic Regression, KNN (test for different values of K), Decision Tree and SVM Classifier**
- 6. Use training and testing to identify the best model**
- 7. Write a summary explaining your results.**

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.style.use('seaborn')
```

```
%matplotlib inline
%pip install graphviz
```

Requirement already satisfied: graphviz in c:\users\apala\anaconda3\lib\site-packages (0.17)

Note: you may need to restart the kernel to use updated packages.

1.- Data import

```
In [5]: from sklearn.datasets import load_breast_cancer
data=load_breast_cancer()
```

2.- Data Visualization and Analysis

2.1 Data Description

The Breast Cancer Dataset is comprised by 569 observations of 30 features and 2 labels. The classes are Malignant (0) and Bening (1). From the 30 features, all of them are numerical and there is no missing date in any of the features observations. Nonetheless, the scale of some features largely differs so we will standardize the features magnitude with the MinMax Scaling Method.

```
In [6]: print(data.DESCR)

.. _breast_cancer_dataset:

Breast cancer wisconsin (diagnostic) dataset
-----

**Data Set Characteristics:**

: Number of Instances: 569

: Number of Attributes: 30 numeric, predictive attributes and the class

: Attribute Information:
  - radius (mean of distances from center to points on the perimeter)
  - texture (standard deviation of gray-scale values)
  - perimeter
  - area
  - smoothness (local variation in radius lengths)
  - compactness (perimeter^2 / area - 1.0)
  - concavity (severity of concave portions of the contour)
  - concave points (number of concave portions of the contour)
  - symmetry
  - fractal dimension ("coastline approximation" - 1)

The mean, standard error, and "worst" or largest (mean of the three
worst/largest values) of these features were computed for each image,
resulting in 30 features. For instance, field 0 is Mean Radius, field
10 is Radius SE, field 20 is Worst Radius.
```

- class:
 - WDBC-Malignant
 - WDBC-Benign

:Summary Statistics:

| | Min | Max |
|-------------------------------------|-------|--------|
| radius (mean): | 6.981 | 28.11 |
| texture (mean): | 9.71 | 39.28 |
| perimeter (mean): | 43.79 | 188.5 |
| area (mean): | 143.5 | 2501.0 |
| smoothness (mean): | 0.053 | 0.163 |
| compactness (mean): | 0.019 | 0.345 |
| concavity (mean): | 0.0 | 0.427 |
| concave points (mean): | 0.0 | 0.201 |
| symmetry (mean): | 0.106 | 0.304 |
| fractal dimension (mean): | 0.05 | 0.097 |
| radius (standard error): | 0.112 | 2.873 |
| texture (standard error): | 0.36 | 4.885 |
| perimeter (standard error): | 0.757 | 21.98 |
| area (standard error): | 6.802 | 542.2 |
| smoothness (standard error): | 0.002 | 0.031 |
| compactness (standard error): | 0.002 | 0.135 |
| concavity (standard error): | 0.0 | 0.396 |
| concave points (standard error): | 0.0 | 0.053 |
| symmetry (standard error): | 0.008 | 0.079 |
| fractal dimension (standard error): | 0.001 | 0.03 |
| radius (worst): | 7.93 | 36.04 |
| texture (worst): | 12.02 | 49.54 |
| perimeter (worst): | 50.41 | 251.2 |
| area (worst): | 185.2 | 4254.0 |
| smoothness (worst): | 0.071 | 0.223 |
| compactness (worst): | 0.027 | 1.058 |
| concavity (worst): | 0.0 | 1.252 |
| concave points (worst): | 0.0 | 0.291 |
| symmetry (worst): | 0.156 | 0.664 |
| fractal dimension (worst): | 0.055 | 0.208 |

:Missing Attribute Values: None

:Class Distribution: 212 - Malignant, 357 - Benign

:Creator: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian

:Donor: Nick Street

:Date: November, 1995

This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.
<https://goo.gl/U2Uwz2>

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

Separating plane described above was obtained using
 Multisurface Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree

Construction Via Linear Programming." Proceedings of the 4th Midwest Artificial Intelligence and Cognitive Science Society, pp. 97-101, 1992], a classification method which uses linear programming to construct a decision tree. Relevant features were selected using an exhaustive search in the space of 1-4 features and 1-3 separating planes.

The actual linear program used to obtain the separating plane in the 3-dimensional space is that described in:

[K. P. Bennett and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets", Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server:

```
ftp ftp.cs.wisc.edu
cd math-prog/cpo-dataset/machine-learn/WDBC/
```

.. topic:: References

- W.N. Street, W.H. Wolberg and O.L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on Electronic Imaging: Science and Technology, volume 1905, pages 861-870, San Jose, CA, 1993.
- O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570-577, July-August 1995.
- W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77 (1994) 163-171.

```
In [7]: x=pd.DataFrame(data.data)
```

```
In [8]: x.shape
```

```
Out[8]: (569, 30)
```

```
In [9]: x.head()
```

```
Out[9]:
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | ... | 20 | 21 |
|---|-------|-------|--------|--------|---------|---------|--------|---------|--------|---------|-----|-------|-------|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | 0.07871 | ... | 25.38 | 17.33 |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | 0.05667 | ... | 24.99 | 23.41 |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | 0.05999 | ... | 23.57 | 25.53 |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | 0.09744 | ... | 14.91 | 26.50 |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | 0.05883 | ... | 22.54 | 16.67 |

5 rows × 30 columns



```
In [10]: data.feature_names
```

```
Out[10]: array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
      'mean smoothness', 'mean compactness', 'mean concavity',
      'mean concave points', 'mean symmetry', 'mean fractal dimension',
      'radius error', 'texture error', 'perimeter error', 'area error',
      'smoothness error', 'compactness error', 'concavity error',
      'concave points error', 'symmetry error',
      'fractal dimension error', 'worst radius', 'worst texture',
      'worst perimeter', 'worst area', 'worst smoothness',
      'worst compactness', 'worst concavity', 'worst concave points',
      'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```


```
In [11]: x.columns=data.feature_names
```

```
In [12]: x.head()
```

```
Out[12]:
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | n fr dimer |
|---|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|------------------|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | 0.0 |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | 0.0 |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | 0.0 |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | 0.0 |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | 0.0 |

5 rows × 30 columns



```
In [13]: y=data.target
```

```
In [14]: print(y)
```

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0
 1 0 0 0 0 0 0 0 0 1 0 1 1 1 1 1 0 0 1 0 0 1 1 1 1 0 1 0 0 1 1 1 0 1 0 0
 1 0 1 0 0 1 1 1 0 0 1 0 0 0 1 1 1 0 1 1 0 0 1 1 1 0 0 1 1 1 0 1 1 0 1 1
 1 1 1 1 1 1 0 0 0 1 0 0 1 1 1 0 0 1 0 1 0 0 1 0 0 1 1 0 1 1 1 1 0 1
 1 1 1 1 1 1 1 1 0 1 1 1 1 0 0 1 0 1 1 0 0 1 1 1 1 0 1 1 0 0 0 1 0
 1 0 1 1 0 1 1 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1 1 0 1 1 0 1 0 0 0 1 1
 1 0 1 1 1 1 1 0 0 1 1 0 1 1 0 0 1 0 1 1 1 1 0 1 1 1 1 0 1 0 0 0 0 0
 0 0 0 0 0 0 0 1 1 1 1 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1
 1 0 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 1
 1 1 0 1 0 1 0 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 1 1 1 1 1 1 1 1 0 0 1 0 0
 0 1 0 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 0 0 1 1 1 1 1 1 1 1 0 1 1 1 1
 1 0 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 1 1 0 1 0 0 1 0 1 1 1 1 1 0 1 1
 0 1 0 1 1 0 1 0 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 1
 1 1 1 1 1 1 0 1 0 1 1 0 1 1 1 1 1 0 0 1 0 1 1 1 1 1 1 0 1 1 0 1 0 0
 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 1 1 1 1 1 1 1 1 0 0 0 0 0 0 1]
```

```
In [15]: print(data.target_names)
```

```
['malignant' 'benign']
```

```
In [16]: types = []
for i in y:
    name = data['target_names'][i]
    types.append(name)
```


```
In [17]: x['types'] = types
```

```
In [18]: x.head()
```

```
Out[18]:
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | n fr dimen |
|---|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|------------------|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | 0.0 |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | 0.0 |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | 0.0 |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | 0.0 |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | 0.0 |

5 rows × 31 columns




```
In [19]: x.tail()
```

```
Out[19]:
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | dim |
|-----|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|-----|
| 564 | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | 0.11590 | 0.24390 | 0.13890 | 0.1726 | 0 |
| 565 | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | 0.10340 | 0.14400 | 0.09791 | 0.1752 | 0 |
| 566 | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | 0.10230 | 0.09251 | 0.05302 | 0.1590 | 0 |
| 567 | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | 0.27700 | 0.35140 | 0.15200 | 0.2397 | 0 |
| 568 | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | 0.04362 | 0.00000 | 0.00000 | 0.1587 | 0 |

5 rows × 31 columns



```
In [20]: x.isna().count()
```

```
Out[20]: mean radius      569
mean texture      569
mean perimeter    569
mean area         569
mean smoothness   569
```

```

mean compactness      569
mean concavity         569
mean concave points   569
mean symmetry         569
mean fractal dimension 569
radius error          569
texture error         569
perimeter error       569
area error            569
smoothness error      569
compactness error     569
concavity error       569
concave points error  569
symmetry error        569
fractal dimension error 569
worst radius          569
worst texture         569
worst perimeter       569
worst area            569
worst smoothness     569
worst compactness     569
worst concavity       569
worst concave points  569
worst symmetry        569
worst fractal dimension 569
types                 569
dtype: int64

```

In [21]:

```
x.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                          569 non-null    float64
1   mean texture                         569 non-null    float64
2   mean perimeter                       569 non-null    float64
3   mean area                           569 non-null    float64
4   mean smoothness                     569 non-null    float64
5   mean compactness                    569 non-null    float64
6   mean concavity                      569 non-null    float64
7   mean concave points                 569 non-null    float64
8   mean symmetry                       569 non-null    float64
9   mean fractal dimension              569 non-null    float64
10  radius error                        569 non-null    float64
11  texture error                       569 non-null    float64
12  perimeter error                     569 non-null    float64
13  area error                          569 non-null    float64
14  smoothness error                    569 non-null    float64
15  compactness error                   569 non-null    float64
16  concavity error                     569 non-null    float64
17  concave points error                569 non-null    float64
18  symmetry error                      569 non-null    float64
19  fractal dimension error             569 non-null    float64
20  worst radius                        569 non-null    float64
21  worst texture                       569 non-null    float64
22  worst perimeter                     569 non-null    float64
23  worst area                          569 non-null    float64
24  worst smoothness                    569 non-null    float64

```

```

25 worst compactness      569 non-null    float64
26 worst concavity        569 non-null    float64
27 worst concave points   569 non-null    float64
28 worst symmetry         569 non-null    float64
29 worst fractal dimension 569 non-null    float64
30 types                  569 non-null    object
dtypes: float64(30), object(1)
memory usage: 137.9+ KB

```

2.1 Feature Standardization

We chose 4 random features to show the difference in the scale of the features and how the MinMax Scaling method helped us to fix that problem by replotting the rescaled features

```


In [23]: import seaborn as sns
x_hist = x[['mean radius', 'mean fractal dimension', 'mean texture', 'mean compactness']]

fig, ax = plt.subplots(ncols=2, nrows=2, figsize=(9, 6))
index=0
ax = ax.flatten()

for col, values in x_hist.items():
    sns.distplot(values, ax=ax[index], kde=False)
    index += 1
plt.tight_layout()

```

C:\Users\apala\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)




```
In [24]: from sklearn.preprocessing import MinMaxScaler
x.iloc[:,0:-1] = MinMaxScaler().fit_transform(x.iloc[:,0:-1])
```

```
In [25]: x.head()
```

```
Out[25]:
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry |
|---|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|
| 0 | 0.521037 | 0.022658 | 0.545989 | 0.363733 | 0.593753 | 0.792037 | 0.703140 | 0.731113 | 0.686364 |
| 1 | 0.643144 | 0.272574 | 0.615783 | 0.501591 | 0.289880 | 0.181768 | 0.203608 | 0.348757 | 0.379798 |
| 2 | 0.601496 | 0.390260 | 0.595743 | 0.449417 | 0.514309 | 0.431017 | 0.462512 | 0.635686 | 0.509596 |
| 3 | 0.210090 | 0.360839 | 0.233501 | 0.102906 | 0.811321 | 0.811361 | 0.565604 | 0.522863 | 0.776263 |
| 4 | 0.629893 | 0.156578 | 0.630986 | 0.489290 | 0.430351 | 0.347893 | 0.463918 | 0.518390 | 0.378283 |

5 rows × 31 columns




```
In [26]: x_hist = x[['mean radius', 'mean fractal dimension', 'mean texture', 'mean compactness']]

fig, ax = plt.subplots(ncols=2, nrows=2, figsize=(9, 6))
index=0
ax = ax.flatten()

for col, values in x_hist.items():
    sns.distplot(values, ax=ax[index], kde=False)
    index += 1
plt.tight_layout()
```

C:\Users\apala\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
In [27]: #from sklearn.preprocessing import StandardScaler
#x.iloc[:,0:-1] = StandardScaler().fit_transform(x.iloc[:,0:-1])
#x["mean radius"].hist()
```

```
In [28]: #x.head()
```


3.- Explanatory Power Analysis

3.1 Correlations Analysis and P-values

As part of the data analysis, we will look for highly correlated features as having multicollinearity implies that there might be some regressors not adding much information to the model. We will try to identify those variables that seem not to have high explanatory power and to create a second sample set without them. For the remaining of the work we will do each analysis with two samples to compare the results

```
In [29]: ## Plot a correlation matrix to see the correlations between (i) features x Label and (
x['target']=y
import seaborn as sns
corr = x.corr()
plt.figure(figsize=(25,12.5))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

```
Out[29]: <AxesSubplot:>
```



From the correlation matrix is easy to see that some variables are highly correlated and others have very low correlation regarding the target variable:

- 1.- Mean Perimeter, mean area, worst radius, worst perimeter and worst area are highly correlated with mean radius
- 2.- Worst texture is highly correlated with mean texture
- 3.- Mean concavity is highly correlated with mean concave points
- 4.- Worst concave points is highly correlated with mean concave points
- 5.- Perimeter error and area error are highly correlated with radius error
- 6.- Perimeter error is highly correlated with area error
- 7.- Worst perimeter and worst area are highly correlated with worst radius
- 8.- Worst area is highly correlated with worst perimeter
- 9.- Mean fractal dimension, texture error, smoothness error, symmetry error, fractal dimension error, have very low correlation with target.

Let us see how the correlation matrix shows after removing the following features:

- Mean concave points, mean texture, perimeter error, area error, worst perimeter, worst area, mean perimeter, mean area, worst radius, worst perimeter and worst area, mean fractal dimension, texture error, smoothness error, symmetry error, fractal dimension error

```
In [30]: x1=x.drop(["mean concave points", "mean texture", "perimeter error", "area error", "worst perimeter", "worst area", "mean perimeter", "mean area", "worst radius", "worst perimeter and worst area", "mean fractal dimension", "texture error", "smoothness error", "symmetry error", "fractal dimension error"])
x1.shape
```


```
Out[30]: (569, 18)
```

```
In [31]:
```

```
corr = x1.corr()
plt.figure(figsize=(20,10))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

Out[31]:

<AxesSubplot:>



We also compute the p-values of the regressors. Normally, a p-value higher than 5% (the common significance level) means that the regressor coefficient is not statistically different from zero. In other words, you could erase that regressor. From the analysis below, it seems that "mean compactness", "concavity error", "worst area" and "worst radius" are the variables with the highest explanatory power.

In [32]:

```
import statsmodels.api as sm
xt=x.drop(["types", "target"],axis=1)
xt2= sm.add_constant(xt)
est = sm.OLS(y, xt2,hasconstant=True)
est2 = est.fit()
print(est2.summary(alpha=0.05))
```

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.774
Model:                  OLS    Adj. R-squared:       0.762
Method:                 Least Squares    F-statistic:        61.53
Date:                   Sun, 17 Oct 2021    Prob (F-statistic):  6.05e-153
Time:                   08:36:11    Log-Likelihood:     29.650
No. Observations:       569    AIC:                2.699
Df Residuals:           538    BIC:                137.4
Df Model:               30
Covariance Type:        nonrobust
=====
```

```
=====
coef    std err          t    P>|t|    [0.025    0.9
75]
```

| | | | | | | |
|-------------------------|---------|-------|--------|-------|---------|-----|
| ----- | | | | | | |
| --- | | | | | | |
| const | 1.7323 | 0.104 | 16.580 | 0.000 | 1.527 | 1. |
| 937 | | | | | | |
| mean radius | 4.6013 | 3.666 | 1.255 | 0.210 | -2.600 | 11. |
| 803 | | | | | | |
| mean texture | -0.1344 | 0.235 | -0.572 | 0.567 | -0.596 | 0. |
| 327 | | | | | | |
| mean perimeter | -3.4354 | 3.632 | -0.946 | 0.345 | -10.570 | 3. |
| 699 | | | | | | |
| mean area | -0.7493 | 1.238 | -0.605 | 0.545 | -3.182 | 1. |
| 683 | | | | | | |
| mean smoothness | -0.0094 | 0.223 | -0.042 | 0.967 | -0.448 | 0. |
| 430 | | | | | | |
| mean compactness | 1.3765 | 0.435 | 3.166 | 0.002 | 0.522 | 2. |
| 231 | | | | | | |
| mean concavity | -0.5967 | 0.446 | -1.337 | 0.182 | -1.474 | 0. |
| 280 | | | | | | |
| mean concave points | -0.4309 | 0.398 | -1.082 | 0.280 | -1.213 | 0. |
| 351 | | | | | | |
| mean symmetry | -0.0203 | 0.147 | -0.138 | 0.890 | -0.309 | 0. |
| 269 | | | | | | |
| mean fractal dimension | -0.0016 | 0.265 | -0.006 | 0.995 | -0.521 | 0. |
| 518 | | | | | | |
| radius error | -1.2011 | 0.857 | -1.401 | 0.162 | -2.885 | 0. |
| 483 | | | | | | |
| texture error | 0.0306 | 0.167 | 0.183 | 0.855 | -0.297 | 0. |
| 358 | | | | | | |
| perimeter error | 0.4779 | 0.873 | 0.548 | 0.584 | -1.236 | 2. |
| 192 | | | | | | |
| area error | 0.4943 | 0.748 | 0.660 | 0.509 | -0.976 | 1. |
| 964 | | | | | | |
| smoothness error | -0.4664 | 0.195 | -2.393 | 0.017 | -0.849 | -0. |
| 084 | | | | | | |
| compactness error | -0.0086 | 0.289 | -0.030 | 0.976 | -0.576 | 0. |
| 559 | | | | | | |
| concavity error | 1.4119 | 0.515 | 2.741 | 0.006 | 0.400 | 2. |
| 424 | | | | | | |
| concave points error | -0.5579 | 0.288 | -1.938 | 0.053 | -1.123 | 0. |
| 007 | | | | | | |
| symmetry error | -0.1206 | 0.194 | -0.622 | 0.534 | -0.501 | 0. |
| 260 | | | | | | |
| fractal dimension error | 0.2069 | 0.338 | 0.612 | 0.541 | -0.457 | 0. |
| 871 | | | | | | |
| worst radius | -5.4866 | 1.629 | -3.367 | 0.001 | -8.688 | -2. |
| 286 | | | | | | |
| worst texture | -0.2686 | 0.261 | -1.030 | 0.303 | -0.781 | 0. |
| 244 | | | | | | |
| worst perimeter | 0.4889 | 1.192 | 0.410 | 0.682 | -1.852 | 2. |
| 830 | | | | | | |
| worst area | 4.1145 | 1.301 | 3.163 | 0.002 | 1.559 | 6. |
| 669 | | | | | | |
| worst smoothness | -0.0822 | 0.217 | -0.378 | 0.705 | -0.509 | 0. |
| 345 | | | | | | |
| worst compactness | -0.0692 | 0.395 | -0.175 | 0.861 | -0.845 | 0. |
| 706 | | | | | | |
| worst concavity | -0.4773 | 0.336 | -1.419 | 0.156 | -1.138 | 0. |
| 183 | | | | | | |
| worst concave points | -0.1351 | 0.266 | -0.508 | 0.612 | -0.658 | 0. |
| 387 | | | | | | |
| worst symmetry | -0.2825 | 0.251 | -1.126 | 0.260 | -0.775 | 0. |

```

210
worst fractal dimension    -0.6561      0.363      -1.806      0.072      -1.370      0.
058
=====
Omnibus:                  32.654      Durbin-Watson:          1.794
Prob(Omnibus):            0.000      Jarque-Bera (JB):       36.690
Skew:                    -0.603      Prob(JB):              1.08e-08
Kurtosis:                 3.302      Cond. No.               923.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

3.2 Features distribution by class

We will create a distribution plot to see how differentiated are each feature in term of the labels, this will also give us an idea of which features may have higher explanatory power such as mean perimeter. Even though, in the middle of the plot there is no clear difference between classes, this feature is clearly differentiated at the limits of the x-axis

In [33]:

```

plt_idx = 1
plt.figure(figsize=(20, 15))

for index1 in range(0,30):
    f1 = x.columns[index1]
    xp = x[f1]
    plt.subplot(6,5,plt_idx,label=f1)
    plt.scatter(xp,y,c = 'red')
    plt.xlabel(f1,labelpad=-75)
    plt_idx = plt_idx+1

```



4.- Metrics for model evaluation

For the model evaluation we will be using the next metrics:

$$\text{Accuracy} = \frac{TP + TN}{FP + FN + TP + TN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

5.- Model Fitting

We will now fit different models to our data. We will apply each model twice. One for the original data and one for the second set in which we remove some variables. We will fit the model to a train set (70% of total data randomly selected) and then predict values for the rest 30%. The models we will apply are Logistic Regression, K-Nearest Neighbor, Decision Tree and Supported Vector Machine

5.1 Logistic Regression

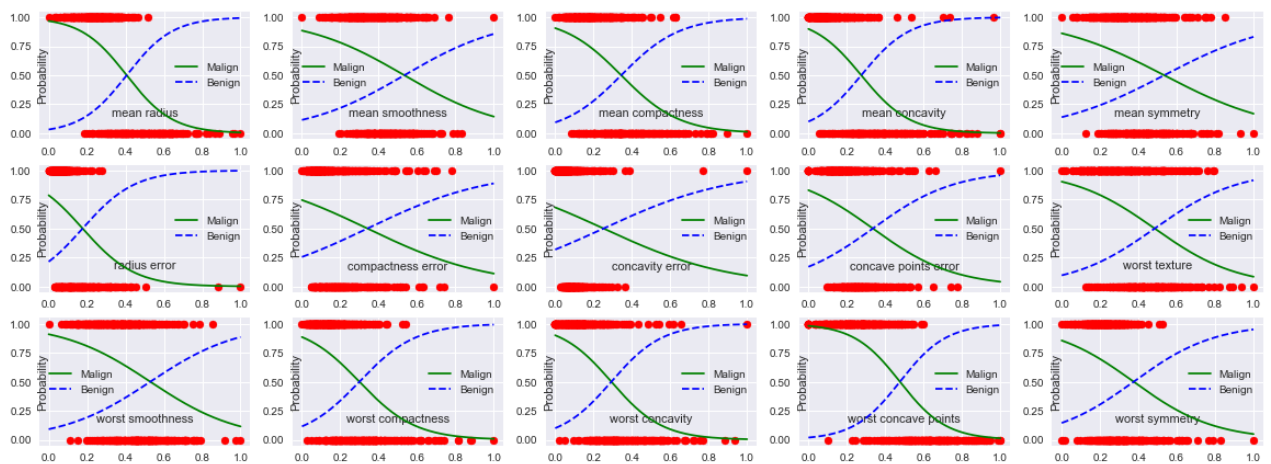
```
In [34]: from sklearn.linear_model import LogisticRegression
```

```

log_reg = LogisticRegression(solver='liblinear')
plt_idx = 1
plt.figure(figsize=(20, 15))

for index1 in range(0,15):
    f1 = x1.columns[index1]
    xp = x1[f1].values.reshape(569,1)
    log_reg.fit(xp,y)
    X_new = np.linspace(xp.min(),xp.max(),1000).reshape(-1,1)
    y_proba = log_reg.predict_proba(X_new)
    plt.subplot(6,5,plt_idx)
    plt.scatter(xp,y,c = 'red')
    plt.plot(X_new,y_proba[:,1],"g-",label="Malign",)
    plt.plot(X_new,y_proba[:,0],"b--",label="Benign")
    plt.xlabel(f1,labelpad=-45,loc='center')
    plt.ylabel('Probability',loc='center',labelpad=-35)
    plt.legend()
    plt_idx = plt_idx+1

```



After cleaning for high correlated variables, its easy to observe that some features do not have high explanatory power when isolated, as the classes ar not clearly differentiated. We will fit the model now with all the features at the same time.

We will define a function generation a confusion matrix, which has on the diagonal the correct prections to show graphically how well the model fit the data. Afterwards, there is a another function that we create so that the data are split between train and test set, the model is fit on the train set and the the metrics described before are applied to evaluate the model

```

In [35]: from sklearn.metrics import plot_confusion_matrix
import matplotlib as mpl

def plot_cm(clf, X, y, labs):

    mpl.rcParams.update({'font.size': 16})
    cm = plot_confusion_matrix(clf, X, y, display_labels=labs,cmap=mpl.cm.Blues);

```

```

In [36]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

```



```

from sklearn.metrics import f1_score

def train(model, z, w):
    x_train, x_test, y_train, y_test = train_test_split(z, w, test_size=.30, random_state=42)
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)

    print('Model Report')
    print('Accuracy: ', np.around(accuracy_score(y_test, y_pred, normalize=True), 4))
    print('Precision: ', np.around(precision_score(y_test, y_pred), 4))
    print('Recall: ', np.around(recall_score(y_test, y_pred), 4))
    print('f1: ', np.around(f1_score(y_test, y_pred), 4))
    plot_cm(model, x_test, y_test, labs=('Bening', 'Malign'))

```

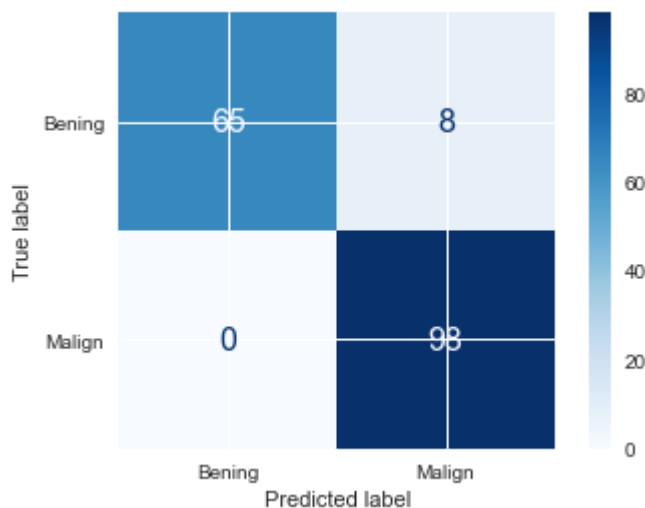
5.1.1 Logistic regression full sample

```

In [38]: xt=x.drop(["types", "target"],axis=1)
         model = LogisticRegression(solver='liblinear')
         train(model, xt, y)

```

Model Report
Accuracy: 0.9532
Precision: 0.9245
Recall: 1.0
f1: 0.9608




5.1.2 Logistic regression subsample

```

In [40]: xt=x1.drop(["types", "target"],axis=1)
         model = LogisticRegression(solver='liblinear')
         train(model, xt, y)

```

Model Report
Accuracy: 0.9415
Precision: 0.9074
Recall: 1.0
f1: 0.9515



The results actually worsen after removing the selected variables, although not by much.


5.2 K-Nearest Neighbor

5.2.1 K-Nearest Neighbor full sample

```
In [41]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [42]: K=3
model = KNeighborsClassifier(K,p=2)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report
Accuracy: 0.9766
Precision: 0.97
Recall: 0.9898
f1: 0.9798




5.2.1.2 K variation analysis

We will see how the accuracy of our prediction changes as we vary the number of neighbors taken to fit the model

```
In [43]: from sklearn.metrics import confusion_matrix
accs = []
ks = []
xt=x.drop(["types", "target"],axis=1)
for k in range(1,101,1):
    x_train, x_test, y_train, y_test = train_test_split(xt, y, test_size=0.30, random_s
    knn = KNeighborsClassifier(n_neighbors = k, p=2)
    knn.fit(x_train, y_train)
    y_pred = knn.predict(x_test)
    cm = confusion_matrix(y_test, y_pred)
    acc = (cm[0][0]+cm[1][1])/len(y_test)
    ks.append(k)
    accs.append(acc)

plt.figure(figsize=(12, 6))
plt.xlabel("K")
plt.ylabel("Accuracy")
plt.plot(ks,accs);
```



5.2.1.2 Check Analysis

Is easy to see that when K=3 we reach the highest accuracy. We will now see how the graph changes when we change the method to measure the distance. The first method used above was the Euclidean method ($p=2$). Now we will use the Manhattan method ($p=1$). Finally, we decided to use the Manhattan method as it reaches the highest accuracy. Now, we will corroborate that accuracy decays as k increases by fitting the model for different ks


```
In [44]: from sklearn.metrics import confusion_matrix
accs = []
ks = []
xt=x.drop(["types", "target"],axis=1)
for k in range(1,101,1):
```

```

x_train, x_test, y_train, y_test = train_test_split(xt, y, test_size=0.30, random_s
knn = KNeighborsClassifier(n_neighbors = k, p=1)
knn.fit(x_train, y_train)
y_pred = knn.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
acc = (cm[0][0]+cm[1][1])/len(y_test)
ks.append(k)
accs.append(acc)

plt.figure(figsize=(12, 6))
plt.xlabel("K")
plt.ylabel("Accuracy")
plt.plot(ks, accs);

```



In [45]:


```

K=3
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)

```

Model Report

Accuracy: 0.9825
Precision: 0.9703
Recall: 1.0
f1: 0.9849



In [47]:

```
K=4
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types","target"],axis=1)
train(model, xt, y)
```


Model Report

Accuracy: 0.9766

Precision: 0.97

Recall: 0.9898

f1: 0.9798



In [48]:

```
K=70
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types","target"],axis=1)
train(model, xt, y)
```


Model Report

Accuracy: 0.924

Precision: 0.8829

Recall: 1.0

f1: 0.9378



5.2.2 K-Nearest Neighbor subsample

In [49]:

```
K=3
model = KNeighborsClassifier(K,p=2)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```

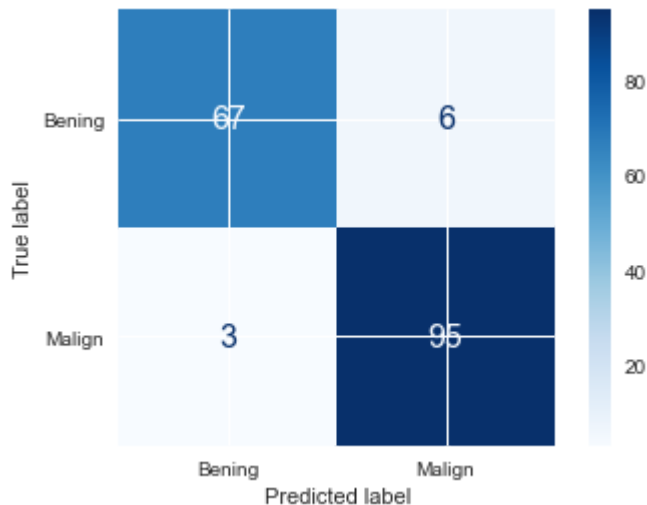
Model Report

Accuracy: 0.9474

Precision: 0.9406

Recall: 0.9694

f1: 0.9548



5.2.2.2 K variation analysis

We will see again how the accuracy of our prediction changes as we vary the number of neighbors taken to fit the model

In [50]:


```
from sklearn.metrics import confusion_matrix
accs = []
ks = []
xt=x1.drop(["types", "target"],axis=1)
for k in range(1,101,1):
    x_train, x_test, y_train, y_test = train_test_split(xt, y, test_size=0.30, random_s
```

```

knn = KNeighborsClassifier(n_neighbors = k, p=2)
knn.fit(x_train, y_train)
y_pred = knn.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
acc = (cm[0][0]+cm[1][1])/len(y_test)
ks.append(k)
accs.append(acc)

plt.figure(figsize=(12, 6))
plt.xlabel("K")
plt.ylabel("Accuracy")
plt.plot(ks,accs);

```



5.2.1.2 Check Analysis

In this case, when K=3 we reach again the highest accuracy. We will repeat the same step before to see how the graph changes when we change the method to measure the distance. In this case, we decided again to use the Manhattan method as it reaches the highest accuracy. Now, we will corroborate that accuracy decays as k increases by fitting the model for different ks


```

In [51]: from sklearn.metrics import confusion_matrix
accs = []
ks = []
xt=x1.drop(["types", "target"],axis=1)
for k in range(1,101,1):
    x_train, x_test, y_train, y_test = train_test_split(xt, y, test_size=0.30, random_s
    knn = KNeighborsClassifier(n_neighbors = k, p=1)
    knn.fit(x_train, y_train)
    y_pred = knn.predict(x_test)
    cm = confusion_matrix(y_test, y_pred)
    acc = (cm[0][0]+cm[1][1])/len(y_test)
    ks.append(k)
    accs.append(acc)

plt.figure(figsize=(12, 6))
plt.xlabel("K")

```

```
plt.ylabel("Accuracy")
plt.plot(ks, accs);
```



In [52]:

```
K=1
model = KNeighborsClassifier(K,p=1)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```


Model Report

Accuracy: 0.9649

Precision: 0.96

Recall: 0.9796

f1: 0.9697




In [53]:

```
K=4
model = KNeighborsClassifier(K,p=1)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report

Accuracy: 0.9708

Precision: 0.9604
Recall: 0.9898
f1: 0.9749




In [54]:

```
K=70  
model = KNeighborsClassifier(K,p=1)  
xt=x1.drop(["types", "target"],axis=1)  
train(model, xt, y)
```

Model Report

Accuracy: 0.9298
Precision: 0.8981
Recall: 0.9898
f1: 0.9417



Similarly as the logistic regression case, the metrics worsen after removing the selected variables but not in a significant way

5.3 Decision Tree

We will now fit the Decision Tree model, setting a maximum depth of ramifications.

In [55]:

```
from sklearn.metrics import classification_report  
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import confusion_matrix
```

5.3.1 Decision Tree full sample

As we can see, all the final leafs have gini measures equal to 0 meaning that the model has exactly fitted the data

```
In [57]: model = DecisionTreeClassifier(random_state=235, max_depth=10)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)
```


Model Report

Accuracy: 0.924

Precision: 0.9126

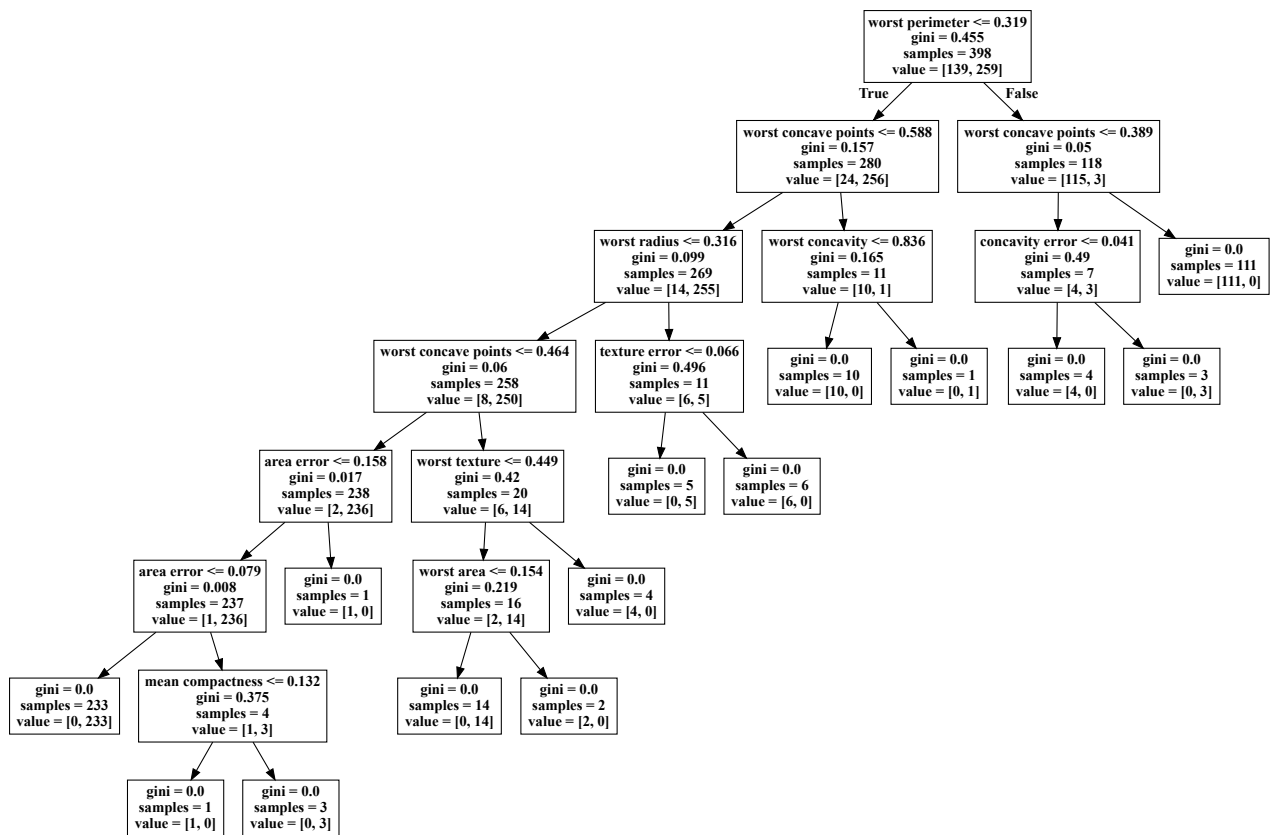
Recall: 0.9592

f1: 0.9353



```
In [58]: from sklearn.tree import export_graphviz
from os import system
from graphviz import Source
from IPython.display import SVG
graph = Source(export_graphviz(model, out_file=None, feature_names = xt.columns))
SVG(graph.pipe(format='svg'))
```

Out[58]:



5.3.2 Decision Tree subsample

In this case, the gini metrics are all equal to zero in the final leafs, so the model fitted the data exactly again. In this case, the subsample actually allowed for an improvement in the performance measures, although they are still lower than the ones obtained in the previous two methods

```
In [59]: model = DecisionTreeClassifier(random_state=235, max_depth=10)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```


Model Report

Accuracy: 0.9415
Precision: 0.9314
Recall: 0.9694
f1: 0.95


```
In [63]: model = SVC(kernel = 'linear', C=1000)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report


Accuracy: 0.9591
Precision: 0.9596
Recall: 0.9694
f1: 0.9645



```
In [64]: model = SVC(kernel = 'rbf', C=1000)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report

Accuracy: 0.9708
Precision: 0.9604
Recall: 0.9898
f1: 0.9749




```
In [65]: model = SVC(kernel = 'poly', C=1000)
xt=x.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report

Accuracy: 0.9591

Precision: 0.9505
Recall: 0.9796
f1: 0.9648




The radial basis kernel yields the best metrics. Let us explore what happens when we dropped the selected variables

5.4.2 Support Vector Machines subsample

```
In [66]: model = SVC(kernel = 'linear', C=1000)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```


Model Report
Accuracy: 0.9474
Precision: 0.9406
Recall: 0.9694
f1: 0.9548



```
In [67]: model = SVC(kernel = 'rbf', C=1000)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report
Accuracy: 0.9649

Precision: 0.96
 Recall: 0.9796
 f1: 0.9697




In [68]:

```
model = SVC(kernel = 'poly', C=1000)
xt=x1.drop(["types", "target"],axis=1)
train(model, xt, y)
```

Model Report

Accuracy: 0.924
 Precision: 0.9208
 Recall: 0.949
 f1: 0.9347



We obtained the same results, the radial basis kernel yields the best result. Also, the analysis is consistent with the logistic and the K-neares neighvor methods, as dropping the selected variables actually worsen the performance

6.- Identify the best method

For selecting the best method, we compily in one talbe the performance metrics for the best case in each method. The best method was the K-Nearest Neighbor with the Manhattan method and K=3 (full sample). The next better ones were the Logistic Regression and the

Support Vector Machines with very little difference between them (full sample). Is interesting to see that the worst method was the one that actually performed better with the subsample.

```
In [71]: L = {'Logistic Regression':[0.9532,0.9245, 1.000, 0.9608], 'KNN':[0.9825,0.9703, 1.0000]
df_results = pd.DataFrame(L, index=['Accuracy', 'Precision', 'Recall', 'F1'])
df_results
```

```
Out[71]:
```

| | Logistic Regression | KNN | Decision Tree | SVM |
|-----------|---------------------|--------|---------------|--------|
| Accuracy | 0.9532 | 0.9825 | 0.9240 | 0.9708 |
| Precision | 0.9245 | 0.9703 | 0.9126 | 0.9604 |
| Recall | 1.0000 | 1.0000 | 0.9592 | 0.9898 |
| F1 | 0.9608 | 0.9849 | 0.9353 | 0.9749 |

7.- Variable significance analysis

We will perform the next analysis with the best model from the table above. We will remove one variable at a time and see how the performance measure of the model changes

```
In [72]: results = pd.DataFrame(index=['Accuracy', 'Precision', 'Recall', 'F1'])
def train_2(model, x, y, name='Feature'):
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=.30, random_state=42)
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)

    stats = [np.around(accuracy_score(y_test, y_pred),4),np.around(precision_score(y_test, y_pred),4),
             np.around(recall_score(y_test, y_pred),4), np.around(f1_score(y_test, y_pred),4)]
    results[name] = stats
```

```
In [74]: K=3
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types","target"],axis=1)
i = 0
for i in range(0, len(data.feature_names)):
    train_2(model, xt.drop(data.feature_names[i],axis=1), y, name = data.feature_names[i])
    i = i + 1


results.sort_values(by="Accuracy",axis=1)
```

```
Out[74]:
```

| | mean texture | worst texture | concave points error | mean radius | worst concavity | worst compactness | worst perimeter | worst radius | concavity error | co |
|-----------|--------------|---------------|----------------------|-------------|-----------------|-------------------|-----------------|--------------|-----------------|----|
| Accuracy | 0.9708 | 0.9708 | 0.9708 | 0.9766 | 0.9766 | 0.9766 | 0.9766 | 0.9766 | 0.9766 | |
| Precision | 0.9604 | 0.9697 | 0.9604 | 0.9700 | 0.9608 | 0.9608 | 0.9700 | 0.9700 | 0.9608 | |
| Recall | 0.9898 | 0.9796 | 0.9898 | 0.9898 | 1.0000 | 1.0000 | 0.9898 | 0.9898 | 1.0000 | |

| | mean texture | worst texture | concave points error | mean radius | worst concavity | worst compactness | worst perimeter | worst radius | concavity error | coi |
|----|-----------------|------------------|----------------------------|----------------|--------------------|----------------------|--------------------|-----------------|--------------------|-----|
| F1 | 0.9749 | 0.9746 | 0.9749 | 0.9798 | 0.9800 | 0.9800 | 0.9798 | 0.9798 | 0.9800 | |

4 rows × 30 columns



The less significant measures are "mean symmetry", "worst area", "mean concave points", "fractal dimension error", "mean compactness", "mean smoothness", "symmetry error", "worst concave points", "mean fractal dimension" and "worst smoothness" as the performance measures are the highest when removing them. The name of the column in the table above indicates which feature has been removed. We will fit the model again now removing those variables

```
In [75]: results=[]
results = pd.DataFrame(index=['Accuracy', 'Precision', 'Recall', 'F1'])
K=4
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types","target","mean symmetry", "worst area", "mean concave points", "frac
names=list(xt.columns)
i = 0
for i in range(0, len(names)):
    train_2(model, xt.drop(names[i],axis=1), y, name = names[i])
    i = i + 1

results.sort_values(by="Accuracy",axis=1)
```

Out[75]:

| | worst texture | worst symmetry | worst fractal dimension | mean texture | mean concavity | radius error | perimeter error | concave points error | compactness error |
|-----------|------------------|-------------------|-------------------------------|-----------------|-------------------|-----------------|--------------------|----------------------------|----------------------|
| Accuracy | 0.9591 | 0.9649 | 0.9708 | 0.9708 | 0.9708 | 0.9708 | 0.9708 | 0.9708 | 0.9708 |
| Precision | 0.9596 | 0.9792 | 0.9794 | 0.9794 | 0.9697 | 0.9697 | 0.9697 | 0.9794 | 0.9794 |
| Recall | 0.9694 | 0.9592 | 0.9694 | 0.9694 | 0.9796 | 0.9796 | 0.9796 | 0.9694 | 0.9694 |
| F1 | 0.9645 | 0.9691 | 0.9744 | 0.9744 | 0.9746 | 0.9746 | 0.9746 | 0.9744 | 0.9744 |



Is interesting how if we remove "Worst Perimeter", we get a better performance measures that without removing any variable

```
In [76]: results=[]
results = pd.DataFrame(index=['Accuracy', 'Precision', 'Recall', 'F1'])
K=4
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types","target","mean symmetry", "worst area", "mean concave points", "frac
names=list(xt.columns)
i = 0
for i in range(0, len(names)):
    train_2(model, xt.drop(names[i],axis=1), y, name = names[i])
    i = i + 1
```

```
results.sort_values(by="Accuracy",axis=1)
```

Out[76]:

| | worst texture | worst symmetry | mean texture | worst fractal dimension | worst compactness | worst concavity | mean concavity | radius error | perimete erro |
|-----------|------------------|-------------------|-----------------|-------------------------------|----------------------|--------------------|-------------------|-----------------|------------------|
| Accuracy | 0.9649 | 0.9649 | 0.9708 | 0.9766 | 0.9766 | 0.9825 | 0.9825 | 0.9825 | 0.9825 |
| Precision | 0.9600 | 0.9694 | 0.9794 | 0.9796 | 0.9796 | 0.9703 | 0.9798 | 0.9703 | 0.9703 |
| Recall | 0.9796 | 0.9694 | 0.9694 | 0.9796 | 0.9796 | 1.0000 | 0.9898 | 1.0000 | 1.0000 |
| F1 | 0.9697 | 0.9694 | 0.9744 | 0.9796 | 0.9796 | 0.9849 | 0.9848 | 0.9849 | 0.9849 |

Finally, we will repeat the analysis by removing the variables we selected from the correlation analysis

In [77]:

```
results=[]
results = pd.DataFrame(index=['Accuracy', 'Precision', 'Recall', 'F1'])
K=3
model = KNeighborsClassifier(K,p=1)
xt=x1.drop(["types", "target"],axis=1)
names=list(xt.columns)
i = 0
for i in range(0, len(names)):
    train_2(model, xt.drop(names[i],axis=1), y, name = names[i])
    i = i + 1

results.sort_values(by="Accuracy",axis=1)
```

Out[77]:

| | mean radius | worst smoothness | mean smoothness | mean symmetry | worst texture | worst concavity | worst concave points | mean compactness | conc |
|-----------|----------------|---------------------|--------------------|------------------|------------------|--------------------|----------------------------|---------------------|------|
| Accuracy | 0.9474 | 0.9474 | 0.9532 | 0.9532 | 0.9532 | 0.9532 | 0.9532 | 0.9591 | 0 |
| Precision | 0.9320 | 0.9238 | 0.9327 | 0.9327 | 0.9412 | 0.9327 | 0.9327 | 0.9417 | 0 |
| Recall | 0.9796 | 0.9898 | 0.9898 | 0.9898 | 0.9796 | 0.9898 | 0.9898 | 0.9898 | 0 |
| F1 | 0.9552 | 0.9557 | 0.9604 | 0.9604 | 0.9600 | 0.9604 | 0.9604 | 0.9652 | 0 |

The results are better when we remove the variables selected in the last feature selection analysis

In [78]:


```
K=4
model = KNeighborsClassifier(K,p=1)
xt=x.drop(["types", "target", "mean symmetry", "worst area", "mean concave points", "frac
train(model, xt, y)
```

Model Report

Accuracy: 0.9883

Precision: 0.98

Recall: 1.0
f1: 0.9899



8.- Conclusion

Even though, all classifiers methods actually performed very well, the K-Nearest Neighbor seems to be the most effective. Nevertheless, given thi high number of variables, the model may have been overfitted. We discard this possibility as we removed several features at a time and the explanatory power remained quite high as reflected by the performance measures yielded when fitting the model on the subsample data. We can conclude that some features, when highly correlated to other variables, do not add much to our classifiers methods, so the most efficient choice is to delete them. Finally, given the performance measures obtained we could rely without doubt into our classifiers to do predicitons regarding the Breast Cancer. A larger sampel would be helpful to have more solid basis.

In [79]: [#aLejandro.palacios@edhec.com](mailto:aLejandro.palacios@edhec.com)